



🌐 [montblanc-project.eu](https://montblanc-project.eu) | @MontBlanc\_EU

## gem5 SVE Hands-On

**Richard Cooper**  
**Arm Research**

**#HiPEAC21**



This project has received funding from the European Union's Horizon 2020 research and innovation program under grant agreement n° 779877

- **Building a gem5 environment (covered this morning).**
- **A very simple example: Saxpy**
  - Compiling the Saxpy example with SVE vectorization for gem5
  - Instrumenting the code for gem5 (m5ops)
  - Running the example in gem5 SE mode
  - Looking at the gem5 output
- **A more realistic example: HACC**
  - Compiling the HACC Example
  - Running HACC in gem5 SE mode

# Prelude: Building a gem5 environment

- **Building gem5 was covered in this morning's session.**
  - A pre-built Docker image is also available on the gem5 website.
  - [https://www.gem5.org/documentation/general\\_docs/building](https://www.gem5.org/documentation/general_docs/building)
- **Quick Reminder (for Ubuntu 20.04).**

```
$ mkdir -p /home/gem5-user/hipec21 ; cd /home/gem5-user/hipec21
$ sudo apt install build-essential git m4 scons zlib1g zlib1g-dev libprotobuf-dev \
    protobuf-compiler libprotoc-dev libgoogle-perftools-dev \
    python3-dev python3-six python-is-python3 libboost-all-dev \
    pkg-config
$ git clone https://gem5.googlesource.com/public/gem5
$ cd gem5
$ scons -j$(nproc) build/ARM/gem5.opt
$ cd ..
$ export GEM5_PATH="/home/gem5-user/hipec21/gem5"
```

- **We will also need to install the aarch64 GCC cross-compiler tools.**

```
$ sudo apt install gcc-aarch64-linux-gnu g++-aarch64-linux-gnu binutils-aarch64-linux-gnu
$
```

# Saxpy example program

```
1 // Saxpy Example for Mont-Blanc Workshop at HiPEAC 2021
2 // Copyright (c) 2020-2021 Arm Limited
3 // All rights reserved.
4 //
5
6 #include <stdlib.h>
7 #include <stdio.h>
8 #include <math.h>
9 #include <time.h>
10
11
12
13 void __attribute__((noinline))
14 saxpy(float * restrict x, float * restrict y, float a, size_t n)
15 {
16     for (size_t i = 0; i < n; ++i)
17     {
18         y[i] = a * x[i] + y[i];
19     }
20 }
21
22 int main(int argc, char * argv[])
23 {
24     if ( argc != 2 ) {
25         fprintf(stderr, "Usage: %s num_elements\n", argv[0]);
26         exit(1);
27     }
28 }
```

```
28     const size_t N = (size_t)atoi(argv[1]);
29     if (N == 0) {
30         fprintf(stderr, "Usage: %s num_elements\n", argv[0]);
31         exit(1);
32     } else {
33         printf("Running saxpy on %ld elements\n", N);
34     }
35
36     float * xs = (float*)malloc(N * sizeof(float));
37     float * ys = (float*)malloc(N * sizeof(float));
38
39     const float a = (float)rand() / (float)RAND_MAX;
40
41     for (size_t i = 0; i < N; ++i) {
42         xs[i] = (float)rand() / (float)RAND_MAX;
43         ys[i] = (float)rand() / (float)RAND_MAX;
44     }
45
46     clock_t start = clock();
47
48     saxpy(xs, ys, a, N);
49
50     clock_t end = clock();
51
52     printf("Elapsed time: %fs\n",
53           ((float)(end - start)) / CLOCKS_PER_SEC);
54
55     free(xs);
56     free(ys);
57
58     exit(0);
59 }
```

- **m5ops are special opcodes that can be inserted into your workload to control the gem5 simulator (e.g. dump statistics, generate checkpoints, etc).**
  - They are encoded in the unused space of the target ISA.
  - [https://www.gem5.org/documentation/general\\_docs/m5ops/](https://www.gem5.org/documentation/general_docs/m5ops/)
- **gem5 provides a C wrapper and library around these instructions for convenience:**
  - Include: `#include "gem5/m5ops.h"`
  - Compile: `-I${GEM5_PATH}/include`
  - Link: `-L${GEM5_PATH}/util/m5/build/aarch64/out -lm5`
- **gem5 also provides the m5 command-line tool which can be called from the console in a full-system simulation. e.g.**
  - `m5 checkpoint`
  - `m5 dump_stats`

# Building the m5ops library

- The m5ops library is built separately as part of the m5 tool.

```
$ cd ${GEM5_PATH}
$ cd util/m5
$ scons build/aarch64/out/m5
$ ls build/aarch64/out
libm5.a  m5
$
```

This can be any of the supported target ISAs:  
x86, arm, thumb, sparc, aarch64

- When compiling your workload, link to `libm5.a` and include `gem5/m5ops.h` to use m5ops in your program.

- `CCFLAGS += -I$(GEM5_PATH)/include`
- `LDFLAGS += -L$(GEM5_PATH)/util/m5/build/aarch64/out`
- `LDFLAGS += -lm5`

# Some useful m5ops

## → Simulation control

- `void m5_exit(uint64_t ns_delay);`
- `void m5_debug_break(void);`
- `void m5_switch_cpu(void);`

We will use `m5_reset_stats(0,0)` and `m5_dump_stats(0,0)` in this example to generate statistics for our region of interest.

## → Statistics generation

- `void m5_reset_stats(uint64_t ns_delay, uint64_t ns_period);`
- `void m5_dump_stats(uint64_t ns_delay, uint64_t ns_period);`
- `void m5_dump_reset_stats(uint64_t ns_delay, uint64_t ns_period);`

## → Checkpoint generation

- `void m5_checkpoint(uint64_t ns_delay, uint64_t ns_period);`

## → Workload delimiters

- `void m5_work_begin(uint64_t workid, uint64_t threadid);`
- `void m5_work_end(uint64_t workid, uint64_t threadid);`

## → And many more...

# Saxpy example program

```
1 // Saxpy Example for Mont-Blanc Workshop at HiPEAC 2021
2 // Copyright (c) 2020-2021 Arm Limited
3 // All rights reserved.
4 //
5
6 #include <stdlib.h>
7 #include <stdio.h>
8 #include <math.h>
9 #include <time.h>
10
11 #include "gem5/m5ops.h"
12
13 void __attribute__((noinline))
14 saxpy(float * restrict x, float * restrict y, float a, size_t n)
15 {
16     for (size_t i = 0; i < n; ++i)
17     {
18         y[i] = a * x[i] + y[i];
19     }
20 }
21
22 int main(int argc, char * argv[])
23 {
24     if ( argc != 2 ) {
25         fprintf(stderr, "Usage: %s num_elements\n", argv[0]);
26         exit(1);
27     }
```

Include m5ops.h from  
the gem5 source tree.

```
28     const size_t N = (size_t)atoi(argv[1]);
29     if (N == 0) {
30         fprintf(stderr, "Usage: %s num_elements\n", argv[0]);
31         exit(1);
32     } else {
33         printf("Running saxpy on %ld elements\n", N);
34     }
35
36     float * xs = (float*)malloc(N * sizeof(float));
37     float * ys = (float*)malloc(N * sizeof(float));
38
39     const float a = (float)rand() / (float)RAND_MAX;
40
41     for (size_t i = 0; i < N; ++i) {
42         xs[i] = (float)rand() / (float)RAND_MAX;
43         ys[i] = (float)rand() / (float)RAND_MAX;
44     }
45
46     clock_t start = clock();
47     m5_reset_stats(0,0);
48     saxpy(xs, ys, a, N);
49     m5_dump_stats(0,0);
50     clock_t end = clock();
51
52     printf("Elapsed time: %fs\n",
53           ((float)(end - start)) / CLOCKS_PER_SEC);
54
55     free(xs);
56     free(ys);
57
58     exit(0);
59 }
```

m5ops allow you to  
annotate your source  
code with special  
instructions for the  
simulator.



# Compiling the Saxpy example

- **GNU gcc can auto-vectorize C code for SVE**
- **Cross compile using the default aarch64-linux-gnu-gcc on Ubuntu 20.04 LTS**
  - On Ubuntu 18.04 LTS, install and compile with aarch64-linux-gnu-gcc-8
- **Use the following compiler switches to enable SVE auto-vectorization:**
  - No vectorization: `CFLAGS += -march=armv8-a+nosimd+nosve -Ofast`
  - Enable SVE vectorization: `CFLAGS += -march=armv8-a+sve -Ofast`
- **For gem5 Syscall Emulation (SE) mode, also link statically:**
  - `LDFLAGS += -static`

# An example Makefile

→ **A simple Makefile to build unoptimized and vectorized versions of the Saxpy example:**

```
# Saxpy Example Makefile for Mont-Blanc Workshop at HiPEAC 2021
# Copyright (c) 2021 Arm Limited
# All rights reserved.
#

# For Ubuntu 20.04 LTS
CC = aarch64-linux-gnu-gcc

# For Ubuntu 18.04 LTS
#CC = aarch64-linux-gnu-gcc-8

GEM5_PATH = ../gem5

CFLAGS = -I$(GEM5_PATH)/include -Wall -Werror -Ofast
CFLAGS_NOOPT = -march=armv8-a+nosimd+nosve
CFLAGS_SVE = -march=armv8-a+sve
LDFLAGS = -static -L$(GEM5_PATH)/util/m5/build/aarch64/out -lm5

.PHONY: all clean

all: saxpy-noopt saxpy-sve

saxpy-noopt: saxpy.c
    $(CC) $(CFLAGS) $(CFLAGS_NOOPT) -o $@ $< $(LDFLAGS)

saxpy-sve: saxpy.c
    $(CC) $(CFLAGS) $(CFLAGS_SVE) -o $@ $< $(LDFLAGS)

clean:
    rm -f saxpy-noopt
    rm -f saxpy-sve
```

# Hands on...

```
gem5-user@ws1:~/hipec21$
```

# Looking at the generated code

→ We can verify that gcc was able to vectorize the saxpy function by looking at the disassembly using objdump.

```
$ aarch64-linux-gnu-gcc \
-march=armv8.4-a+nosimd+nosve -Ofast \
-I../gem5/include \
-o saxpy-noopt saxpy.c \
-static -L../gem5/util/m5/build/aarch64/out -lm5
$ aarch64-linux-gnu-objdump -d saxpy-noopt | less
...

0000000000400810 <saxpy>:
400810:      b4001122      cbz     x2, 400834 <saxpy+0x24>
400814:      d2800003      mov     x3, #0x0
400818:      bc637802      ldr     s2, [x0, x3, lsl #2]
40081c:      bc637821      ldr     s1, [x1, x3, lsl #2]
400820:      1f000441      fmadd   s1, s2, s0, s1
400824:      bc237821      str     s1, [x1, x3, lsl #2]
400828:      91000463      add     x3, x3, #0x1
40082c:      eb03005f      cmp     x2, x3
400830:      54ffff41      b.ne    400818 <saxpy+0x8>
400834:      d65f03c0      ret

...
```

```
$ aarch64-linux-gnu-gcc \
-march=armv8.4-a+sve -Ofast \
-I../gem5/include \
-o saxpy-sve saxpy.c \
-static -L../gem5/util/m5/build/aarch64/out -lm5
$ aarch64-linux-gnu-objdump -d saxpy-sve | less
...

0000000000400810 <saxpy>:
400810:      b4001a2        cbz     x2, 400844 <saxpy+0x34>
400814:      d2800003      mov     x3, #0x0                // #0
400818:      05242002      mov     z2.s, s0
40081c:      25a21fe0      whilelo p0.s, xzr, x2
400820:      2598e3e1      ptrue   p1.s
400824:      d503201f      nop
400828:      a5434020      ld1w    {z0.s}, p0/z, [x1, x3, lsl #2]
40082c:      a5434001      ld1w    {z1.s}, p0/z, [x0, x3, lsl #2]
400830:      65a10440      fmla    z0.s, p1/m, z2.s, z1.s
400834:      e5434020      st1w    {z0.s}, p0, [x1, x3, lsl #2]
400838:      04b0e3e3      incw    x3
40083c:      25a21c60      whilelo p0.s, x3, x2
400840:      54ffff41      b.ne    400828 <saxpy+0x18>    // b.any
400844:      d65f03c0      ret

...
```

# Looking at the generated code

→ We can verify that gcc was able to vectorize the saxpy function by looking at the disassembly using objdump.

```
$ aarch64-linux-gnu-gcc \
-march=armv8.4-a+nosimd \
-I../gem5/include \
-o saxpy-neon saxpy.c \
-I../gem5/util/ \
-o saxpy-noopt saxpy.c \
-static -L../gem5/util/ \
$ aarch64-linux-gnu-objdump -d
```

```
000000000400810 <saxpy>:
400810: b400122      400810: b400122
400814: d280003      400814: d280003
400818: bc637802     400818: bc637802
40081c: bc637821     40081c: bc637821
400820: 1f000441     400820: 1f000441
400824: bc237821     400824: bc237821
400828: 91000463     400828: 91000463
40082c: eb03005f     40082c: eb03005f
400830: 54ffff41     400830: 54ffff41
400834: d65f03c0     400834: d65f03c0
```

```
$ aarch64-linux-gnu-gcc \
-march=armv8.4-a+nosimd -Ofast \
-I../gem5/include \
-o saxpy-neon saxpy.c \
-I../gem5/util/ \
-o saxpy-noopt saxpy.c \
-static -L../gem5/build/aarch64/out -lm5
$ aarch64-linux-gnu-objdump -d saxpy-sve | less
...
000000000400810 <saxpy>:
400810: b4000482      cbz    x2, 4008a0 <saxpy+0x90>
400814: d1000443      sub    x3, x2, #0x1
400818: f100007f      cmp    x3, #0x2
40081c: 540004a9      b.ls   4008a0 <saxpy+0x94>
400820: d342fc44      lsr    x4, x2, #2
400824: d2800003      mov    x3, #0x0
400828: 4e040403      dup    v3.4s, v0.s[0]
40082c: 037ce084      lsl    v4.4s, v4.4s, #4
400830: 3ce36821      ldr     q1, [x1, x3]
400834: 3ce36802      ldr     q2, [x0, x3]
400838: 4e23cc41      fmla   v1.4s, v2.4s, v3.4s
40083c: 3ca36821      str     q1, [x1, x3]
400840: 91000463      add    x3, x3, #0x10
400844: eb04007f      cmp    x3, x4
400848: 54ffff41      b.ne   400830 <saxpy+0x20>
40084c: f240045f      tst     x2, #0x3
400850: 927ef443      and    x3, x2, #0xffffffffffffc
400854: 54000260      b.eq   4008a0 <saxpy+0x90>
400858: bc637802      ldr     s2, [x0, x3, lsl #2]
40085c: 91000464      add    x4, x3, #0x1
400860: bc637821      ldr     s1, [x1, x3, lsl #2]
400864: eb04005f      cmp    x2, x4
400868: 1f000441      fmaddd s1, s2, s0, s1
40086c: bc237821      str     s1, [x1, x3, lsl #2]
400870: 54000189      b.ls   4008a0 <saxpy+0x90>
400874: bc647802      ldr     s2, [x0, x4, lsl #2]
400878: 91000863      add    x3, x3, #0x2
40087c: bc647821      ldr     s1, [x1, x4, lsl #2]
400880: eb03005f      cmp    x2, x3
400884: 1f020401      fmaddd s1, s0, s2, s1
400888: bc247821      str     s1, [x1, x4, lsl #2]
40088c: 540000a9      b.ls   4008a0 <saxpy+0x90>
400890: bc637802      ldr     s2, [x0, x3, lsl #2]
400894: bc637821      ldr     s1, [x1, x3, lsl #2]
400898: 1f020400      fmaddd s0, s0, s2, s1
40089c: bc237820      str     s0, [x1, x3, lsl #2]
4008a0: d65f03c0      ret
4008a4: d2800003      mov    x3, #
4008a8: 27ffffec      b      400850 <saxpy+0x48>
...
```

```
ldr    q1, [x1, x3]
ldr    q2, [x0, x3]
fmla   v1.4s, v2.4s, v3.4s
str     q1, [x1, x3]
add     x3, x3, #0x10
cmp     x3, x4
b.ne    400830 <saxpy+0x20>
```

```
...
build/aarch64/out -lm5
py-sve | less
...
x2, 400844 <saxpy+0x34>
x3, #0x0 // #0
z2.s, s0
p0.s, xzr, x2
true p1.s
p
lw {z0.s}, p0/z, [x1, x3, lsl #2]
lw {z1.s}, p0/z, [x0, x3, lsl #2]
la z0.s, p1/m, z2.s, z1.s
lw {z0.s}, p0, [x1, x3, lsl #2]
cw x3
leilo p0.s, x3, x2
ne 400828 <saxpy+0x18> // b.any
t
```

# Running the Saxpy example in gem5 SE mode

- To run a SVE program in gem5 SE mode, the only special thing we need to do is set the SVE vector length.
- Do this using the `--param` switch of `se.py`
  - `--param` sets a parameter of the `SimObjects` in the simulator
  - In this case we are setting the `sve_vl_se` parameter of all the `Isa` objects under all `Cpu` objects under the `System` object.
  - In gem5 the `sve_vl_se` is an integer multiple of 128-bits. In the example below, `sve_vl_se = 4` means the simulation will use an SVE vector length of 512-bits.

```
$ cd ${GEM5_PATH}
$ ./build/ARM/gem5.opt configs/example/se.py \
    ...
    --param 'system.cpu[:].isa[:].sve_vl_se = 4' \
    ...
```

This is a special syntax supported by the `--param` switch. It means 'all items in the isa collection'.

# A simple system configuration

## → A simple run-script for the Saxpy example:

```
#!/bin/bash

GEM5_PATH=../gem5

sve_vl=4

${GEM5_PATH}/build/ARM/gem5.opt \
    ${GEM5_PATH}/configs/example/se.py \
    --cpu-type MinorCPU \
    --mem-type SimpleMemory \
    --cmd saxpy-sve --options 6000 \
    --caches --l2cache \
    --l1i_size=64kB --l1i_assoc=4 \
    --l1d_size=64kB --l1d_assoc=4 \
    --l2_size=256kB --l2_assoc=4 \
    --mem-size=1GB \
    --cacheline_size=128 \
    --param "system.cpu[:].isa[:].sve_vl_se = ${sve_vl}"
```

# Hands on...

```
gem5-user@ws1:~/hipec21/saxpy$ make
aarch64-linux-gnu-gcc -I../gem5/include -Wall -Werror -Ofast -march=armv8-a+nosimd+nosve -o saxpy-noopt saxpy.c -static -L../gem5/util/m5/build/aarch64/out -lm5
aarch64-linux-gnu-gcc -I../gem5/include -Wall -Werror -Ofast -march=armv8-a+sve -o saxpy-sve saxpy.c -static -L../gem5/util/m5/build/aarch64/out -lm5
gem5-user@ws1:~/hipec21/saxpy$ ll saxpy-*
-rwxrwxrwx 1 gem5-user gem5-user 605968 Jan 14 20:44 saxpy-noopt*
-rwxrwxrwx 1 gem5-user gem5-user 605968 Jan 14 20:44 saxpy-sve*
gem5-user@ws1:~/hipec21/saxpy$
```



# Looking at the gem5 Statistics

- ➔ The gem5 output is stored in ./m5out by default, or a directory specified by the user using the --outdir/-d flag.

[http://learning.gem5.org/book/part1/gem5\\_stats.html](http://learning.gem5.org/book/part1/gem5_stats.html)

- config.ini, config.json: A record of the system configuration.
- simerr, simout: The simulation output, if run with the -r and -e flags.
- stats.txt: The output statistics of the simulation.

```
----- Begin Simulation Statistics -----  
final_tick                624412500  
(restored from checkpoints and never reset)  
host_inst_rate            17533140  
host_mem_usage            264700  
host_op_rate              21349632  
host_seconds              0.05  
host_tick_rate            66255762  
sim_freq                  1000000000000  
sim_insts                  842001  
sim_ops                   1035265  
sim_seconds               0.000003  
sim_ticks                 3219500  
....
```

```
# Number of  
# Simulated  
# Number of  
# Simulated  
# Real time  
# Simulated  
# Frequency of simulated ticks  
# Number of instructions simulated  
# Number of ops (including micro ops) simulated  
# Number of seconds simulated  
# Number of ticks simulated
```

- stats.txt may contain multiple statistics blocks:
- One block is generated for each call to m5\_dump\_stats() or m5\_dump\_reset\_stats().
- A final block is generated when the simulation ends.

# Looking at the gem5 Statistics - Simulation Time

----- Begin Simulation Statistics -----

final_tick	624412500	# Number of ticks from beginning of simulation (restored from...
host_inst_rate	17533140	# Simulator instruction rate (inst/s)
host_mem_usage	264700	# Number of bytes of host memory used
host_op_rate	21349632	# Simulator op (including micro ops) rate (op/s)
host_seconds	0.05	# Real time elapsed on the host
host_tick_rate	66255762	# Simulator tick rate (ticks/s)
sim_freq	1000000000000	# Frequency of simulated ticks
sim_insts	842001	# Number of instructions simulated
sim_ops	1035265	# Number of ops (including micro ops) simulated
<u>sim_seconds</u>	<u>0.000003</u>	# Number of <u>seconds simulated</u>
<u>sim_ticks</u>	<u>3219500</u>	# Number of <u>ticks simulated</u>
system.cpu.committedInsts	2642	# Number of instructions committed
system.cpu.committedOps	2642	# Number of ops (including micro ops) committed
system.cpu.cpi	2.437169	# CPI: cycles per instruction
system.cpu.discardedOps	17	# Number of ops (including micro ops) which were discarded...
system.cpu.idleCycles	2	# Total number of cycles that the object has spent stopped
system.cpu.ipc	0.410312	# IPC: instructions per cycle
<u>system.cpu.numCycles</u>	<u>6439</u>	# number of <u>cpu cycles simulated</u>
system.cpu.numFetchSuspend	0	# Number of times Execute suspended instruction fetching
system.cpu.numWorkItemsCompleted	0	# number of work items this cpu completed
system.cpu.numWorkItemsStarted	0	# number of work items this cpu started
....		

# Looking at the gem5 Statistics - Instruction Counts

----- Begin Simulation Statistics -----

....

system.cpu.op_class_0::No_OpClass	0	0.00%	0.00% # Class of committed instruction
system.cpu.op_class_0::IntAlu	388	14.69%	14.69% # Class of committed instruction
system.cpu.op_class_0::IntMult	0	0.00%	14.69% # Class of committed instruction
system.cpu.op_class_0::IntDiv	0	0.00%	14.69% # Class of committed instruction
...			
system.cpu.op_class_0::FloatMult	0	0.00%	14.69% # Class of committed instruction
<u>system.cpu.op_class_0::FloatMultAcc</u>	<u>0</u>	<u>0.00%</u>	<u>14.69% # Class of committed instruction</u>
system.cpu.op_class_0::FloatDiv	0	0.00%	14.69% # Class of committed instruction
system.cpu.op_class_0::FloatMisc	1	0.04%	14.72% # Class of committed instruction
...			
system.cpu.op_class_0::SimdFloatMult	0	0.00%	43.19% # Class of committed instruction
<u>system.cpu.op_class_0::SimdFloatMultAcc</u>	<u>375</u>	<u>14.19%</u>	<u>57.38% # Class of committed instruction</u>
system.cpu.op_class_0::SimdFloatSqrt	0	0.00%	57.38% # Class of committed instruction
...			
system.cpu.op_class_0::SimdShaSigma3	0	0.00%	57.38% # Class of committed instruction
system.cpu.op_class_0::SimdPredAlu	1	0.04%	57.42% # Class of committed instruction
<u>system.cpu.op_class_0::MemRead</u>	<u>750</u>	<u>28.39%</u>	<u>85.81% # Class of committed instruction</u>
<u>system.cpu.op_class_0::MemWrite</u>	<u>375</u>	<u>14.19%</u>	<u>100.00% # Class of committed instruction</u>
system.cpu.op_class_0::FloatMemRead	0	0.00%	100.00% # Class of committed instruction
system.cpu.op_class_0::FloatMemWrite	0	0.00%	100.00% # Class of committed instruction

...

# Hands on...

```
gem5-user@ws1:~/hipec21/saxpy$
```

# Comparing different SVE vector lengths

→ A simple run-script for the Saxpy example:

```
#!/bin/bash

GEM5_PATH=../gem5

for sve_vl in $(seq 1 16) ; do
    outdir=results/saxpy-sve-vl${sve_vl}
    mkdir -p ${outdir}
    ${GEM5_PATH}/build/ARM/gem5.opt \
        --outdir ${outdir} \
        ${GEM5_PATH}/configs/example/se.py \
        --cpu-type MinorCPU \
        --mem-type SimpleMemory \
        --cmd saxpy-sve --options 6000 \
        --caches --l2cache \
        --l1i_size=64kB --l1i_assoc=4 \
        --l1d_size=64kB --l1d_assoc=4 \
        --l2_size=256kB --l2_assoc=4 \
        --mem-size=1GB \
        --cacheline_size=128 \
        --param "system.cpu[:].isa[:].sve_vl_se = ${sve_vl}"
done
```

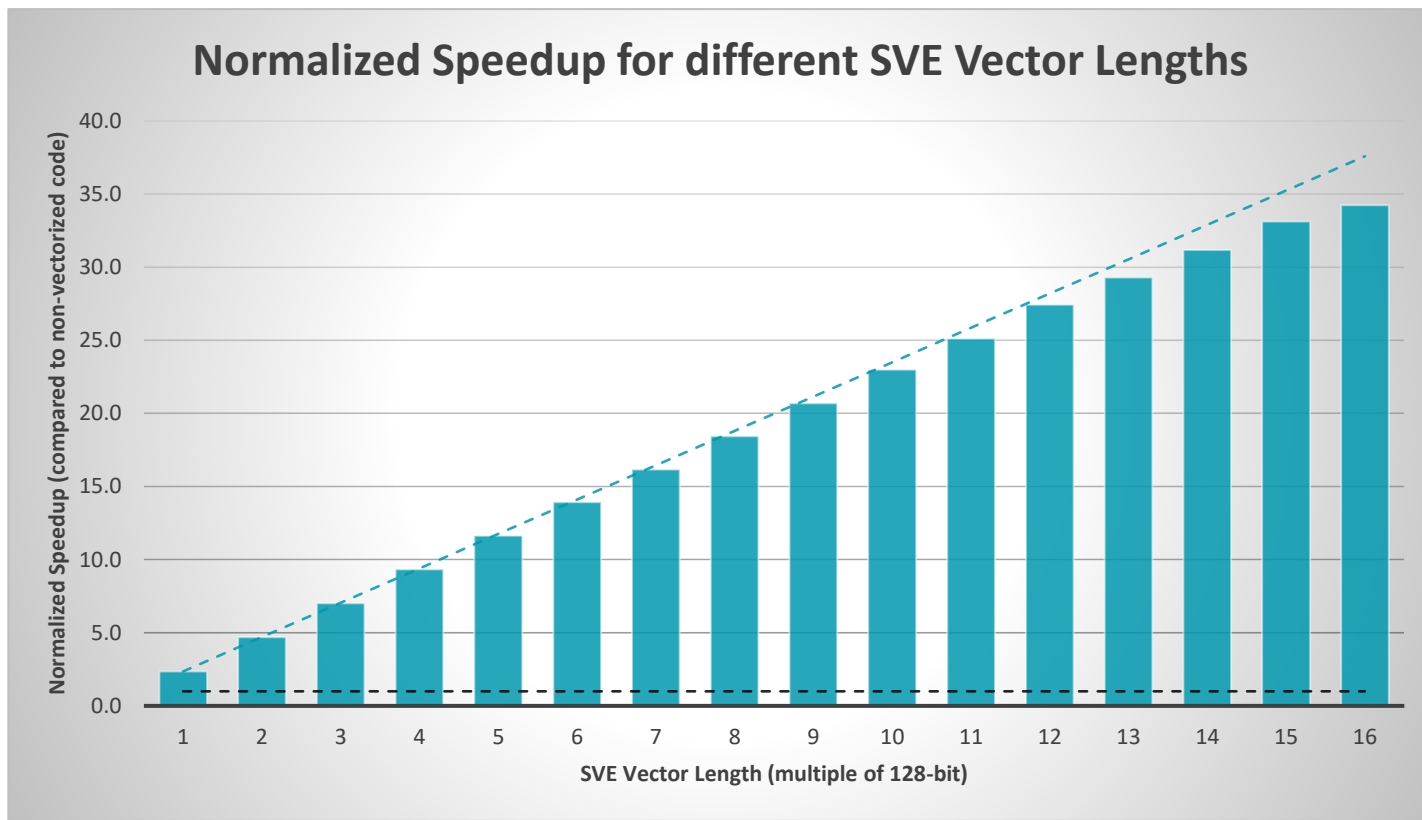
# Comparing different SVE vector lengths

- **Each stats.txt file will have two statistics blocks.**
  - The first contains the statistics for the region between `m5_reset_stats()` and `m5_dump_stats()`.
  - The second contains the stats at the end of the simulation.
- **Extract the relevant statistics from the first block to compare the SVE simulations for different vector lengths.**
  - `sim_seconds`, `sim_ticks`, `system.cpu.numCycles`
  - `system.cpu.op_class_0::`
    - `{FloatMultAcc, SimdFloatMultAcc, MemRead, MemWrite}`
- **The statistics of interest can easily be extracted manually or using your favourite tools (e.g. Python).**

# Comparing different SVE vector lengths

Stats Filename	sim_seconds	sim_ticks	numCycles	FloatMultAcc	SimdFloatMultAcc	MemRead	MemWrite	numCycles <sub>noop</sub> /numCycles
results/saxpy-6000-noopt-vl1/stats.txt	0.0000300	30034000	60068	6000	0	12000	6000	1.00000
results/saxpy-6000-sve-vl1/stats.txt	0.0000130	12782000	25564	0	1500	3000	1500	2.34971
results/saxpy-6000-sve-vl2/stats.txt	0.0000060	6407000	12814	0	750	1500	750	4.68769
results/saxpy-6000-sve-vl3/stats.txt	0.0000040	4282000	8564	0	500	1000	500	7.01401
results/saxpy-6000-sve-vl4/stats.txt	0.0000030	3219500	6439	0	375	750	375	9.32878
results/saxpy-6000-sve-vl5/stats.txt	0.0000030	2582000	5164	0	300	600	300	11.63207
results/saxpy-6000-sve-vl6/stats.txt	0.0000020	2157000	4314	0	250	500	250	13.92397
results/saxpy-6000-sve-vl7/stats.txt	0.0000020	1859500	3719	0	215	430	215	16.15165
results/saxpy-6000-sve-vl8/stats.txt	0.0000020	1630000	3260	0	188	376	188	18.42577
results/saxpy-6000-sve-vl9/stats.txt	0.0000010	1451500	2903	0	167	334	167	20.69170
results/saxpy-6000-sve-vl10/stats.txt	0.0000010	1307000	2614	0	150	300	150	22.97934
results/saxpy-6000-sve-vl11/stats.txt	0.0000010	1196500	2393	0	137	274	137	25.10155
results/saxpy-6000-sve-vl12/stats.txt	0.0000010	1094500	2189	0	125	250	125	27.44084
results/saxpy-6000-sve-vl13/stats.txt	0.0000010	1025000	2050	0	116	232	116	29.30146
results/saxpy-6000-sve-vl14/stats.txt	0.0000010	963500	1927	0	108	216	108	31.17177
results/saxpy-6000-sve-vl15/stats.txt	0.0000010	907000	1814	0	100	200	100	33.11356
results/saxpy-6000-sve-vl16/stats.txt	0.0000010	877500	1755	0	94	188	94	34.22678

# Comparing different SVE vector lengths





# What is gem5 Simulating

## → These results look too good to be true...

- Very simple workload and region of interest.
- Small working set.
- I chose a slightly unrealistic system configuration so the workload would not be memory bound.

## → What does gem5 simulate out of the box?

- Not any existing CPU...
  - Generic in-order and out-of-order microarchitecture.
  - Operation latencies are not tuned for a specific CPU.
  - Surrounding memory system is very configurable.

## → Always know what you are simulating

- Sometimes the out-of-the box behaviour is sufficient - e.g. to compare relative performance.
- For other investigations it may be necessary to tune the model:
  - Accurate custom CPU models.
  - Tune operation latencies to match a specific CPU.
  - Configure the surrounding system to match a target platform.
- *“The gem5 simulator is a modular platform for computer-system architecture research.”*

```
saxpy(float * restrict x,
      float * restrict y,
      float a,
      size_t n)
{
    for (size_t i = 0; i < n; ++i)
    {
        y[i] = a * x[i] + y[i];
    }
}
```

```
#!/bin/bash
```

```
GEM5_PATH=../gem5
```

```
for sve_vl in $(seq 1 16) ; do
    outdir=results/saxpy-sve-vl${sve_vl}
    mkdir -p ${outdir}
    ${GEM5_PATH}/build/ARM/gem5.opt \
        --outdir ${outdir} \
        ${GEM5_PATH}/configs/example/se.py \
        --cpu-type MinorCPU \
        --mem-type SimpleMemory \
        --cmd saxpy-sve --options 6000 \
        --caches --l2cache \
        --l1i_size=64kB --l1i_assoc=4 \
        --l1d_size=64kB --l1d_assoc=4 \
        --l2_size=256kB --l2_assoc=4 \
        --mem-size=1GB \
        --cacheline_size=128 \
        --param "system.cpu[.].isa[.].sve_vl_se = ${sve_vl}"
done
```

# Where to find the Op Class Names

→ The SVE instructions are defined in:

`gem5/src/arch/arm/isa/insts/sve.isa` & `sve_mem.isa`

e.g.

# FMLA (vectors)

`sveTerInst('fmla', 'Fmla', 'SimdFloatMultAccOp', floatTypes, fmlaCode, PredType.MERGE)`

Opcode

Op Class

→ The Op Class Latencies are defined in the CPU code for each CPU type.

- Tuning the CPU models' Op Class latencies is beyond the scope of this tutorial, but the basic process is to derive a CPU class from one of the base CPU classes and provide a custom Functional Unit Pool.
- See the config scripts in `gem5/configs/common/cores/arm` for examples.

# A more complex example: HACC

- Now let's try to run the HACC example from the Arm SVE Tools Tutorial: <https://gitlab.com/arm-hpc/training/arm-sve-tools>
- This is the HACCKernels Benchmark for Hardware/Hybrid Accelerated Cosmology Code (HACC) - see the README for more details.
- Uses OpenMP and can be SVE vectorized.
- Requires a couple of small changes to the make configuration for cross compilation (see Hands-on).
  - arm-sve-tools/config.mk:34
    - CFLAGS\_OPT = -Ofast -mcpu=native
    - + CFLAGS\_OPT = -Ofast -mcpu=generic -static
  - arm-sve-tools/05\_Apps/01\_HACC/Makefile:54
    - \$(CXX) \$(CXXFLAGS\_REPORT) \$(CXXFLAGS\_OPT) \$(CXXFLAGS\_OPENMP) -o \$\$@ \$\$^
    - + \$(CXX) \$(CXXFLAGS\_REPORT) \$(CXXFLAGS\_OPT) -march=armv8-a+sve \$(CXXFLAGS\_OPENMP) -o \$\$@ \$\$^

# A more complex example: HACC

## → Building the example:

```
$ cd /home/gem5-user/hipeac21
$ git clone https://gitlab.com/arm-hpc/training/arm-sve-tools.git
$ # Edit config.mk here...
$ cd arm-sve-tools/05_Apps/01_HACC
$ make COMPILER=gcc CC=aarch64-linux-gnu-gcc CXX=aarch64-linux-gnu-g++
..... Compiler output .....
$
```

## → Running the example:

```
$ ${GEM5_PATH}/build/ARM/gem5.opt \
    ${GEM5_PATH}/configs/example/se.py \
    --cpu-type DerivO3CPU --num-cpus ${NUM_CPUS} --mem-type SimpleMemory --mem-size=1GB \
    --cmd hacc_gnu_sve.exe --options 4 \
    --caches --l2cache
    --l1i_size=64kB --l1i_assoc=4 \
    --l1d_size=64kB --l1d_assoc=4
    --l2_size=256kB --l2_assoc=4 \
    --cacheline_size=128 --param "system.cpu[:].isa[:].sve_vl_se = 4
```

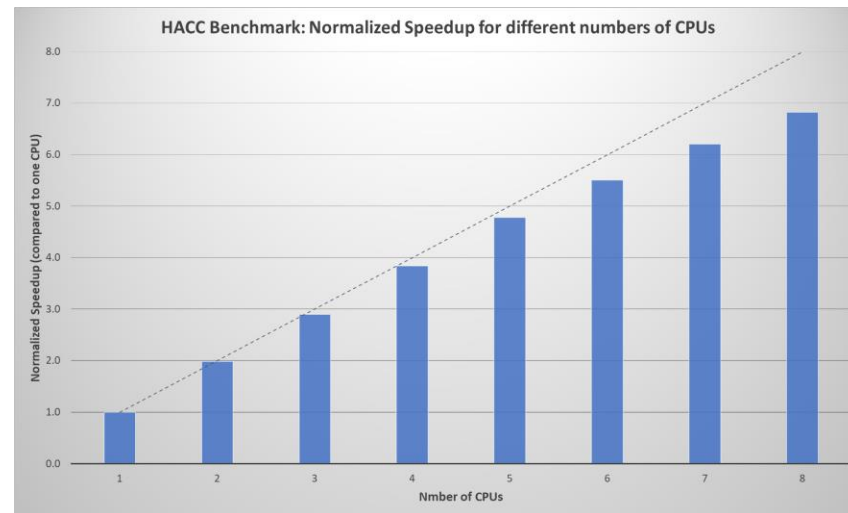
# Hands on...

```
gem5-user@ws1:~/hipec21$ git clone https://gitlab.com/arm-hpc/training/arm-sve-tools.git
Cloning into 'arm-sve-tools'...
remote: Enumerating objects: 1035, done.
remote: Counting objects: 100% (1035/1035), done.
remote: Compressing objects: 100% (513/513), done.
remote: Total 1406 (delta 580), reused 952 (delta 510), pack-reused 371
Receiving objects: 100% (1406/1406), 250.93 MiB | 5.49 MiB/s, done.
Resolving deltas: 100% (744/744), done.
Updating files: 100% (354/354), done.
gem5-user@ws1:~/hipec21$ cp ../run-hacc-simulation.sh arm-sve-tools/05_Apps/01_HACC
gem5-user@ws1:~/hipec21$
```

# HACC Benchmark: Simulated Speedup

Number of CPUs	sim_seconds	sim_ticks	sim_ticks <sub>(1 CPU)</sub> /sim_ticks
1	0.091960	91959629000	1.0000
2	0.046387	46386878000	1.9824
3	0.031744	31743740000	2.8969
4	0.023985	23985125000	3.8340
5	0.019242	19242198000	4.7791
6	0.016710	16709786000	5.5033
7	0.014827	14826835000	6.2022
8	0.013489	13488657000	6.8176

SVE Vector Length = 4



# Running SVE programs in Full-System Mode

- SVE also works in gem5 Full System Mode.
- In `gem5/configs/example/fs.py`, configure the *maximum* SVE vector length using the `--sve-vl` command line parameter (this is different to `se.py` mode).
- Inside Linux, the default SVE Vector Length can be read or set for new processes using the `procfs` interface.
  - Read:  
`cat /proc/sys/abi/sve_default_vector_length`
  - Set:  
`echo ${vl} > /proc/sys/abi/sve_default_vector_length`
  - Note: Linux uses the number of bytes to describe the SVE Vector Length:  
128-bits ⇔ `sve_default_vector_length=16` (Linux) ⇔ `sve_vl=1` (gem5)

SVE Vector Length	gem5 sve_vl	Linux
128 (1 x 128-bits)	1	16
256 (2 x 128-bits)	2	32
384 (3 x 128-bits)	3	48
...	...	...

# Concluding Remarks

- ➔ **Running SVE Programs in gem5 doesn't require any special consideration, apart from configuring the SVE Vector Length.**
- ➔ **Pay attention to the system that is being simulated: is it realistic?**
  - What counts as realistic depends on your research purpose.



# MONT-BLANC 2020

## QUESTION TIME

**More information:**  
[montblanc-project.eu](http://montblanc-project.eu)  
[@MontBlanc\\_EU](https://twitter.com/MontBlanc_EU)